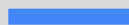


Python Tools to Deploy Your Machine Learning Models Faster 🚀



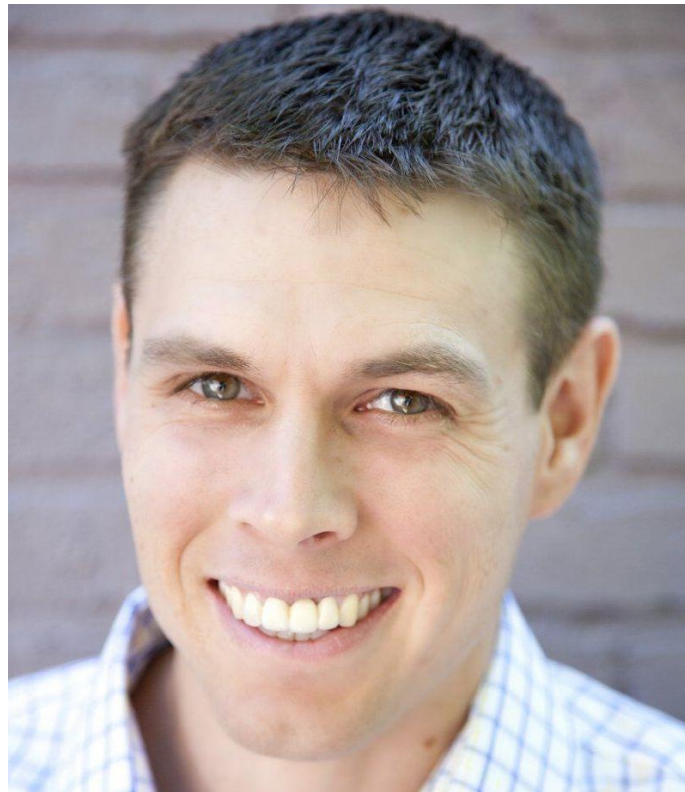
Jeff Hale



Welcome Aboard!

Pilot: Jeff Hale

- [linkedin.com/in/-jeffhale/](https://www.linkedin.com/in/-jeffhale/)
- [twitter@discdiver](https://twitter.com/discdiver)
- jeffhale.medium.com/



Preflight check

What's one thing you're hoping to learn today?

Planes



Flight plans

- Test flights for each plane
 - Hello world
 - Show data
 - Plot
 - Model inference
- Cruising altitude & turbulence (pros & cons)
- Grab your luggage (takeaways)
- Disembark

Flight plan materials

- GitHub repo:

<https://github.com/discdiver/dsdc-deploy-models>

- Slides:

<https://www.slideshare.net/JeffHale6/getting-a-job-in-data-science-what-and-how-to-learn>

Gradio

Ultralight



New, quick to fly, experimental



Gradio Demo 1: Hello world



Gradio #1: Hello world

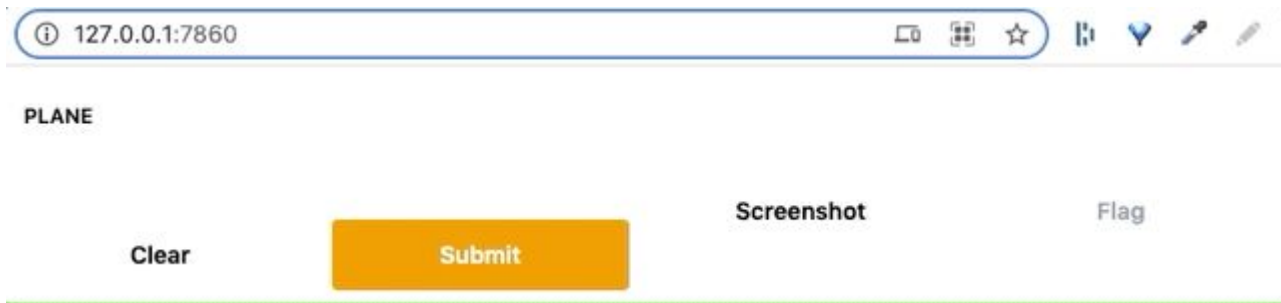
```
import gradio as gr

def hello(plane):
    return f"I'm an ultralight {plane} ✈️"

iface = gr.Interface(
    fn=hello,
    inputs=['text'],
    outputs=['text']
).launch()
```

Gradio #1: Hello world

- `pip install gradio`
- `python gradio_hello.py`





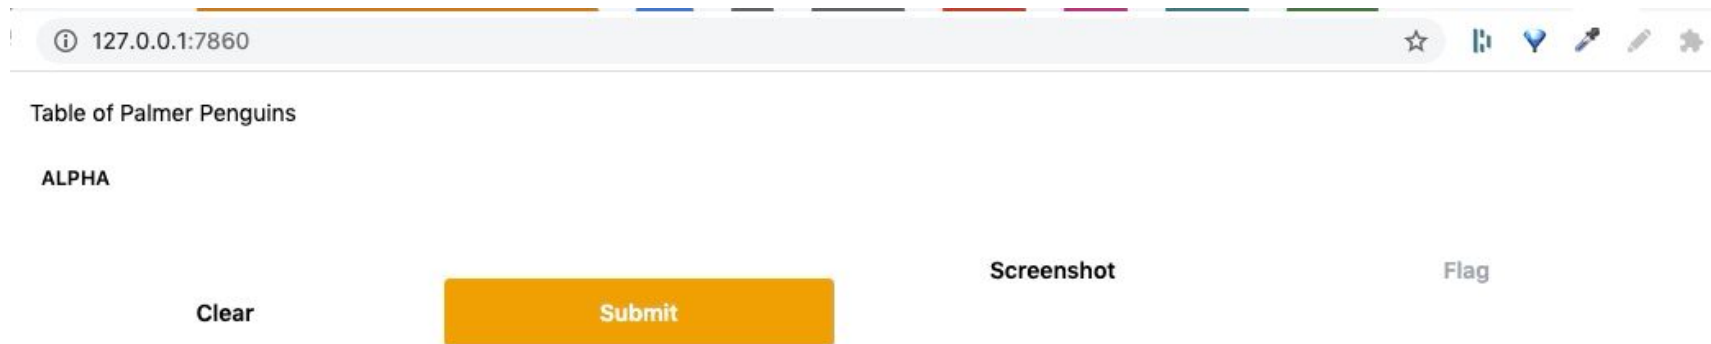
Gradio Demo 2: Show data

Gradio #2: Show me the data!

```
def show_pens(alpha):  
    return pd.read_csv(  
        'https://raw.githubusercontent.com/mwaskom/seaborn-data/master/penguins.csv'  
    )  
  
iface = gr.Interface(  
    fn=show_pens,  
    inputs=['text'],  
    outputs=[gr.outputs.DataFrame()],  
    description="Table of Palmer Penguins"  
).launch(share=True)
```

Gradio #2: Show me the data!

- `pip install gradio pandas seaborn`
- `python gradio_pandas.py`
- Glitch only allows 10 cells to show initially or fails silently as of Nov. 16, 2021 TK



The screenshot shows a web browser window with the address bar displaying '127.0.0.1:7860'. The page content includes a title 'Table of Palmer Penguins', a label 'ALPHA', and a large orange 'Submit' button. There are also 'Clear', 'Screenshot', and 'Flag' buttons. The interface is part of a Gradio application running locally.

127.0.0.1:7860

Table of Palmer Penguins

ALPHA

Clear Submit Screenshot Flag



Gradio Demo 3: Plotting

Gradio #3: Plot it

- Plotly doesn't work as of 2.8.7 (targeted for 2.9) 😞
- You can use Matplotlib as of Gradio 2.8.2

Gradio #3: Plot it

```
def plot_pens(place_holder):  
    """scatter plot penguin chars using matplotlib"""  
    df_pens = pd.read_csv(  
        "https://raw.githubusercontent.com/mwaskom/seaborn-data/master/penguins.csv"  
    )  
    fig = plt.figure()  
    plt.scatter(x=df_pens["bill_length_mm"],  
y=df_pens["bill_depth_mm"])  
    return fig
```

Gradio #3: Plot it

```
iface = gr.Interface(  
    fn=plot_pens,  
    layout="vertical",  
    inputs=["checkbox"],  
    outputs=["plot"],  
    title="Scatterplot of Palmer Penguins",  
    description="Let's talk pens. Click to see a plot.",  
    article="Talk more about Penguins here, shall we?",  
    theme="peach",  
    live=True,  
).launch()
```

Scatterplot of Palmer Penguins

Let's talk pens. Click to see a plot.

place holder

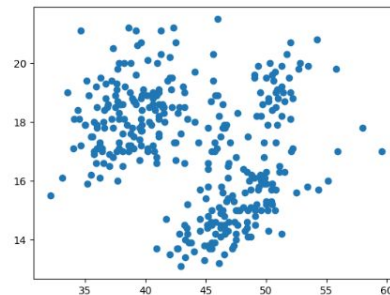


Clear

Submit

Output

0.3s



Flag

Talk more about Penguins here, shall we?

[view the api](#) 🔑 • built with [gradio](#) 📦



Gradio Demo 4: Model inference

Gradio #4: Model inference

```
import gradio as gr  
gr.Interface.load('huggingface/gpt2').launch()
```

gpt2

INPUT

Clear

Submit

Screenshot

Flag

Gradio #4: Model inference - prettier

```
gr.Interface.load(
    "huggingface/gpt2",
    title="Storytelling with GPT2",
    css="""
        body {background: rgb(2,0,36);
            background: linear-gradient(
                180deg,
                rgba(2,0,36,1) 0%,
                rgba(7,51,99,1) 70%,
                rgba(6,3,17,1) 100%);}
        .title {color: white !important;}
    """,
).launch()
```

Gradio #4: Model inference - prettier

Storytelling with GPT2

INPUT

I wish I were a fish

Clear

Submit

OUTPUT

1.68s

I wish I were a fish. I don't know if I've ever done that, but I still am a fish for good measure -- because, by the end of the day, maybe I got one back. That wasn't the case. I

Screenshot

Flag

Gradio Data API - One Click!

Response:

```
{  
  "data": [ Union[str, number] ],  
  "durations": [ float ], // the time taken for the prediction to complete  
  "avg_durations": [ float ] // the average time taken for all predictions so far (used to  
estimate the runtime)  
}
```

Try it (live demo):

Python

cURL

Javascript

```
curl -X POST http://127.0.0.1:7860/api/predict -H 'Content-Type: application/json' -d  
'{"data": ["Hello World"]}'
```


Gradio Pros

- Quick demos for ML 🚀
- Built-in interpretability 🔍
- Auto-docs 📄
- Nice Hugging Face integration 🤗
- Bright future 😎

Gradio Cons

- Rough around the edges - early stage 🗡️
- Not easy to customize elements ✨
- Single page only 1

Gradio - Hugging Face Acquisition



+



Abubakar Abid

Streamlit

Cessna Citation Longitude



Light, quick to takeoff, easy flying



Streamlit Demo 1: Hello world

Streamlit #1: Hello world

```
import streamlit as st  
  
name = "Jeff"  
  
st.title(f"Hello from Streamlit, {name}!")
```

- *pip install streamlit*
- *streamlit run streamlit_hello.py*

Hello from Streamlit!



Streamlit Demo 2: Show data

Streamlit #2: Show data

```
import streamlit as st
import pandas as pd

st.title("Streamlit with pandas")
show = st.checkbox("Show dataframe")
df_pens = pd.read_csv(
    "https://raw.githubusercontent.com/mwaskom/seaborn-data/master/penguins.csv")

if show:
    df_pens
```


Streamlit #2: Show data

- *pip install streamlit pandas seaborn*
- *streamlit run streamlit_pandas.py*

Streamlit #2: Show data

Streamlit with pandas

☐ Show dataframe



Streamlit Demo 3: Plotting

Streamlit #3: Plotting

```
choice = st.radio("Select color", ["species", "island"])
```

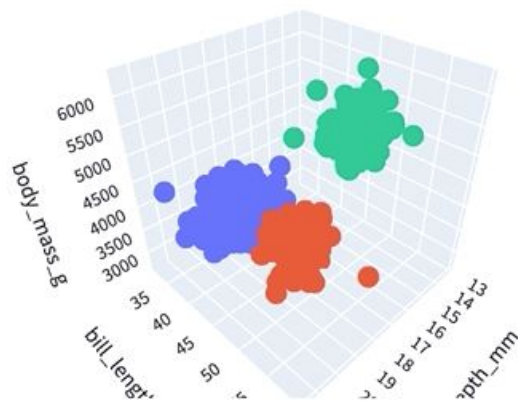
```
fig = px.scatter_3d(  
    data_frame=df_pens,  
    x="bill_depth_mm",  
    y="bill_length_mm",  
    z="body_mass_g",  
    color=choice,  
    title="Penguins in 3D!",  
)  
fig
```

Streamlit #3: Plotting

Select color

- ☒ species
- ☐ island

Penguins in 3D!



species

- Adelie
- Chinstrap
- Gentoo



Streamlit

Streamlit Demo 4: Model Inference

Streamlit #4: Model inference

```
import streamlit as st

from transformers import pipeline

st.header("GPT-2 Stories")

input_text = st.text_area("Enter your text here:")

generator = pipeline("text-generation", model="gpt2")

output = generator(input_text, max_length=100, )

output[0]["generated_text"]
```

Streamlit #4: Model inference

GPT-2 Stories

Enter your text here:

ok friends, let's talk

ok friends, let's talk

Riot - I'm on the board now. Please explain your message.

No one likes news on my channel.

Please put your email in the contact form. - Yes I have a problem

Streamlit #4: Model inference - prettier

GPT-2 Stories

Enter your text here:

ok friends, let's talk

|

ok friends, let's talk

Riot - I'm on the board now. Please explain your message.

No one likes news on my channel.

Please put your email in the contact form. - Yes I have a problem

Streamlit #4: Model inference - prettier

```
st.header("Story time")

st.image("https://cdn.pixabay.com/photo/2017/07/12/19/03/highway-2497900_960_720.jpg")

col1, col2 = st.columns(2)

with col1:

    input_text = st.text_area("Enter your text here:")

    with st.spinner("Generating story..."):

        generator = pipeline("text-generation", model="gpt2")

        if input_text:

            generated_text = generator(input_text, max_length=60)

            st.success("Here's your story:")

            generated_text[0]["generated_text"]

with col2:

    st.header("I'm in another column")
```

Streamlit #4: Model inference - prettier

Story time



Enter your text here:

ok smarty, here we are

I'm in another column

Here's your story:

ok smarty, here we are.

There's also the recent news of IBM's (IBM) (IBM Inc) (IBM) (IBM) (IBM MS) (IBM) smartwatches and the new BlackBerry (R) in this week's Best

Streamlit Serving Options

- Serve from Streamlit's servers for free: bit.ly/st-6plots
- Serve from Hugging Face or Heroku for free
- Pay Streamlit for more/better
- Host elsewhere

Streamlit Pros

- Quick websites for many Python use cases 🦆
- Many intuitive interactive widgets ✅
- Nice hosting options 📺
- Thoughtful docs 📄
- Strong development cadence & team 💪

Streamlit Cons

- Some customizations cumbersome ✨
- Single page only (without hacks) 1

Streamlit Snowflake Acquisition

Recent acquisition by Snowflake



FastAPI

Boeing 737



Commercial grade, fast, smart!



FastAPI Demo 1: Hello world

FastAPI #1: Hello world

```
import uvicorn

from fastapi import FastAPI


app = FastAPI()


@app.get("/")
def home():
    return {"Hello world": "How's it going?"}


if __name__ == "__main__":
    uvicorn.run("fastapi_hello:app")
```

FastAPI #1: Hello world

```
pip install fastapi uvicorn
```

```
python fastapi_hello.py
```

Returns json

```
{"Hello world": "How's it going ?"}
```

FastAPI: Async



FastAPI #1: Hello world

```
@app.get("/not-async")  
  
def home():  
    return dict(zip(range(10), range(10)))
```

10ms

FastAPI #1: Hello world

```
@app.get("/yes-async")
```

```
async def home():
```

```
    return dict(zip(range(10), range(10)))
```

3ms



FastAPI Demo 2: Show data

FastAPI #2: Show me the data!

```
@app.get("/df")  
  
async def pens_data():  
    df_pens = pd.read_csv(  
        "https://raw.githubusercontent.com/mwaskom/seaborn-d  
ata/master/penguins.csv")  
  
    df_no_nans = df_pens.fillna(-1.01)  
  
    return df_no_nans
```


Automatic docs

FastAPI 0.1.0 OAS3
/openapi.json

default

GET / Home

GET /not-async Demo1

GET /yes-async Demo2

GET /df Pens Data

Parameters Cancel

No parameters

Execute Clear

Responses

Curl

```
curl -X 'GET' \
  'http://127.0.0.1:8000/df' \
  -H 'accept: application/json'
```

Request URL

```
http://127.0.0.1:8000/df
```

FastAPI #2: Show me the data!

Get back the whole DataFrame as json. Don't need to convert. 😎

```
{"species":{"0":"Adelie","1":"Adelie"...
```



FastAPI Demo 3: Plotting

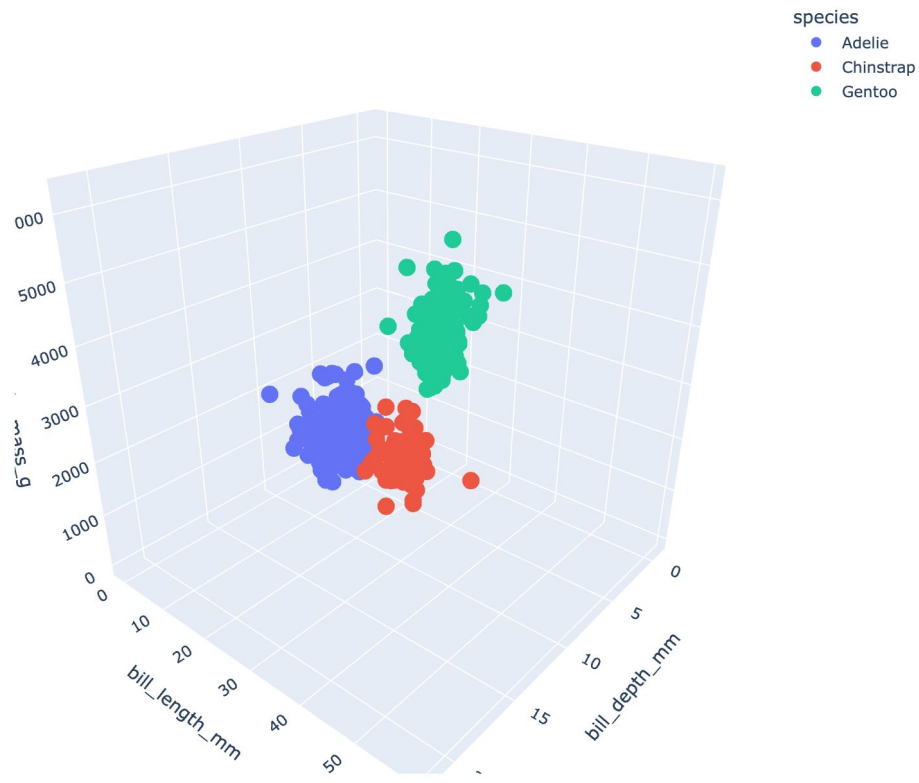
FastAPI #3: Plotting

```
@app.get("/plot")  
  
async def plot() -> HTMLResponse:  
    """return a plotly plot"""  
  
    fig = px.scatter_3d(  
        data_frame=df,  
        x="bill_depth_mm",  
        y="bill_length_mm",  
        z="body_mass_g",  
        color="species",  
        title="Penguins in 3D!",  
    )  
  
    return HTMLResponse(fig.to_html())
```

FastAPI #3: Plotting



Penguins in 3D!












FastAPI Demo 4: Model Inference

FastAPI #4: Model inference

Typing with mypy TK

FastAPI Pros

- Fastest Python API framework - async 
- Automatic API documentation 
- Nice error messages 
- Extensive docs 
- Jinja templating 
- Nice test client 
- SQL Model integration 

FastAPI Pros




Sebastian Ramirez

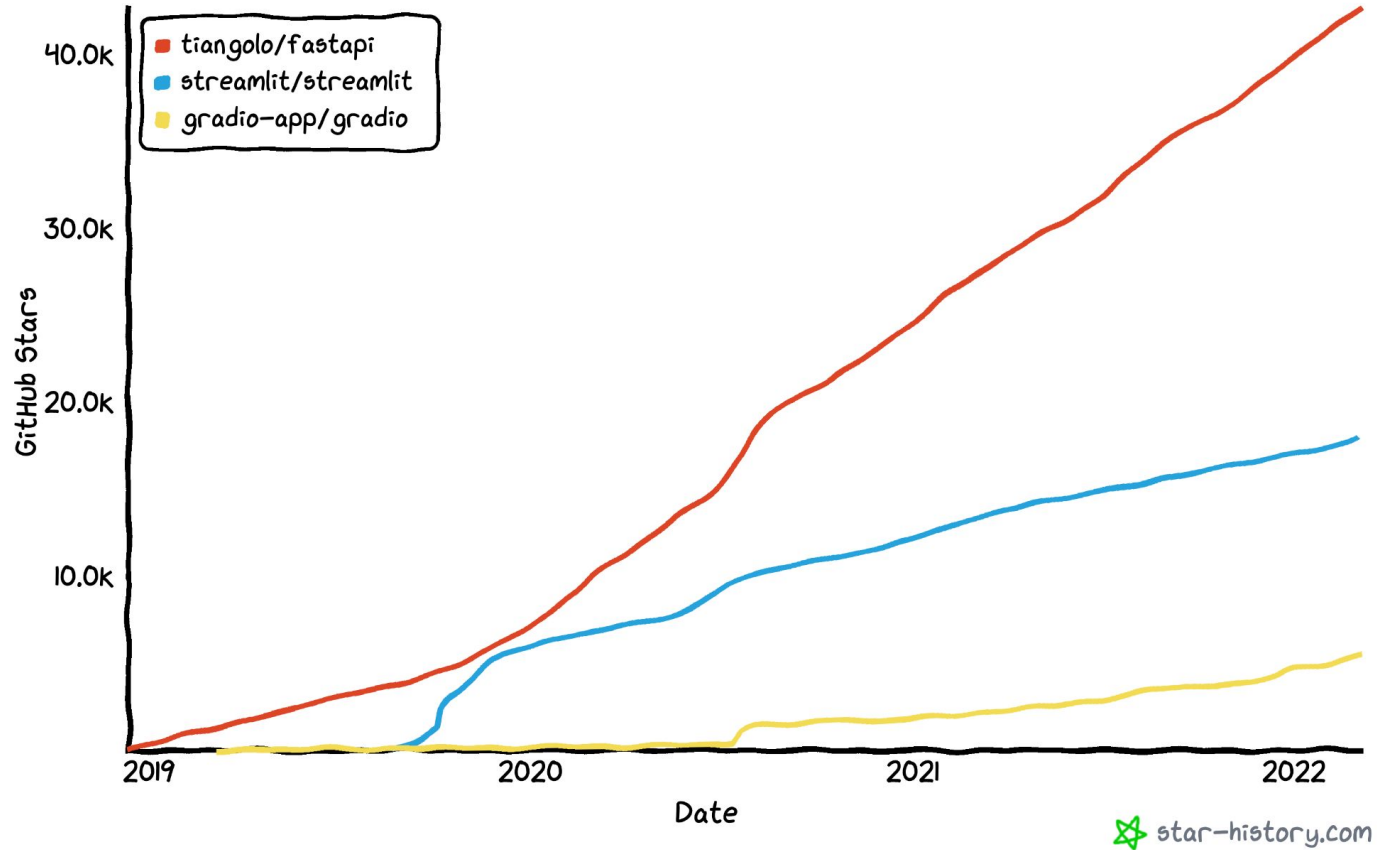
FastAPI Cons

- Reliant on a single maintainer 🤖
- Takes more code for HTML templating than Flask ⌨️

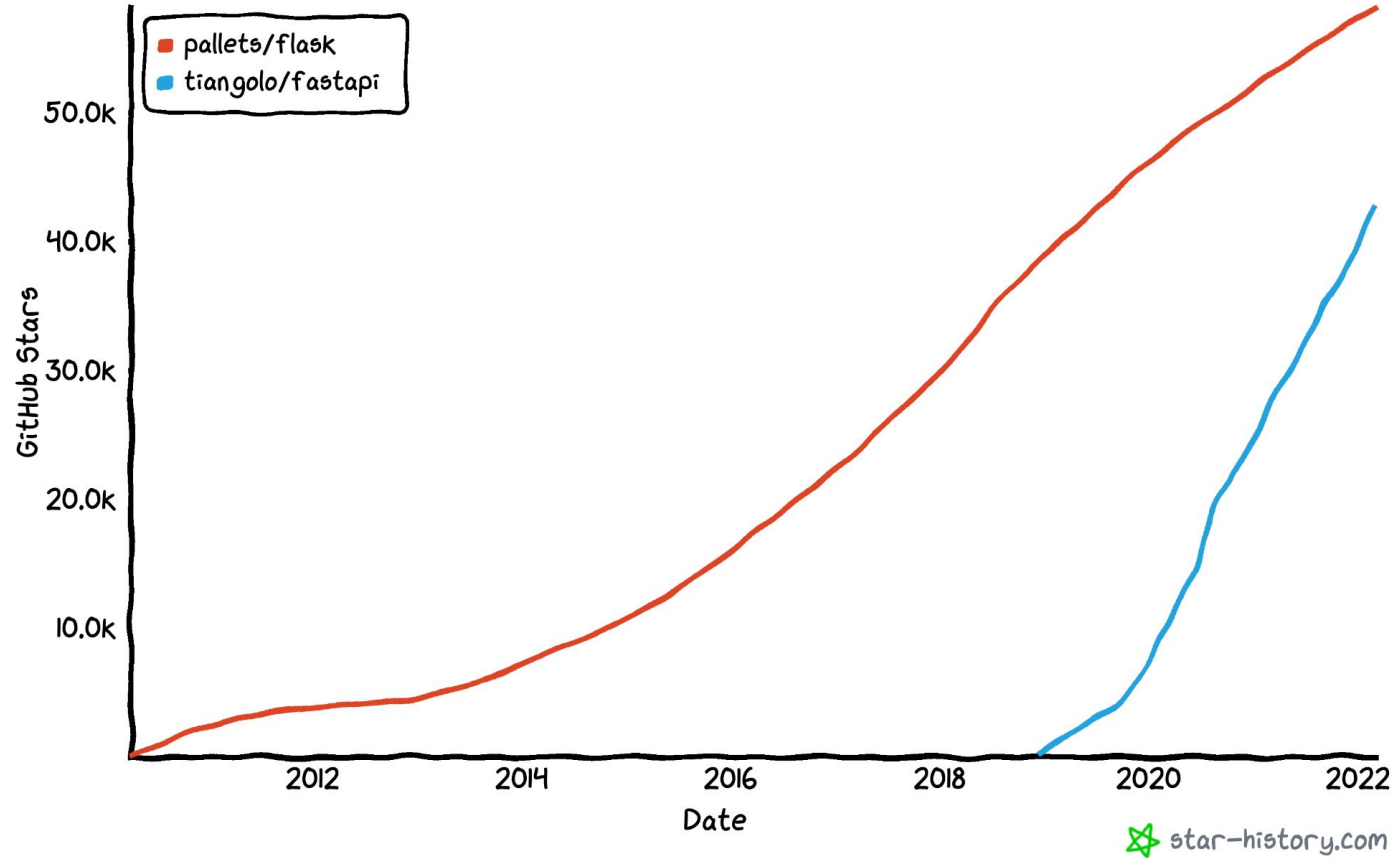
Summin' Up

	Web App	Data API
	Yes	Yes
 Streamlit	Yes	No
 FastAPI	Yes (Jinja templates)	Yes

Star history



Star history



star-history.com

**Grab Your Luggage
(takeaways)**

**Use what you know, unless it
doesn't meet your needs.**

Blank slate?



**Learn what's popular, growing,
and quick to get off the ground.**



Streamlit



**For single-page app that doesn't
need custom styling.**



Gradio for quick 🤗 models for fun.



FastAPI for serving data.

What to learn next?

(Newer to Python)



Streamlit

Disembark

Thank you for flying the deployment skies!

Jeff Hale

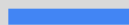
- [linkedin.com/in/-jeffhale/](https://www.linkedin.com/in/-jeffhale/)
- [@discdiver](#)
- jeffhale.medium.com/



Versions used

- Gradio 2.8.7
- Streamlit 1.5.1
- FastAPI 0.75.0

Python Tools to Deploy Your Machine Learning Models Faster 🚀



Jeff Hale

