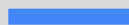


Python Tools to Deploy Your Machine Learning Models Faster 🚀



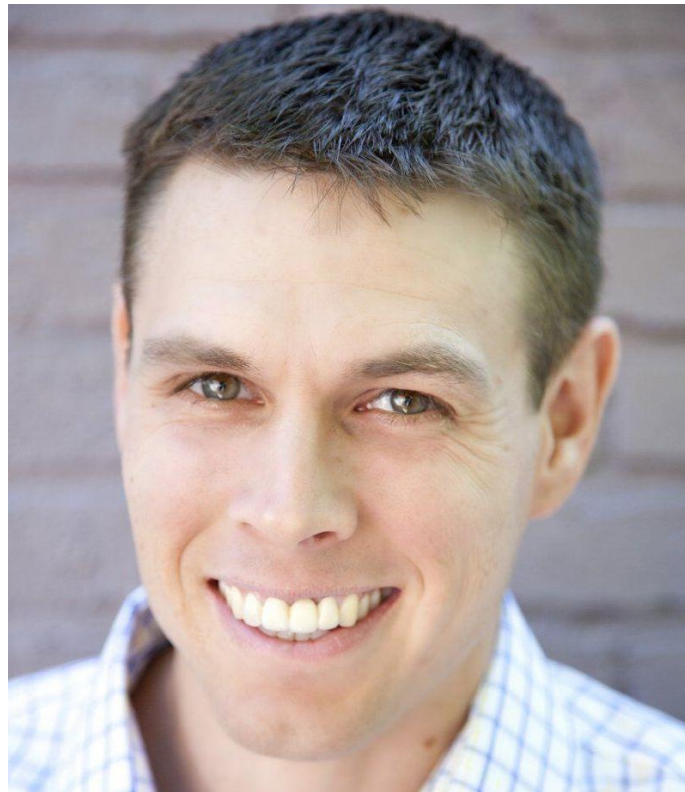
Jeff Hale



Welcome Aboard!

Pilot: Jeff Hale

- [linkedin.com/in/-jeffhale/](https://www.linkedin.com/in/-jeffhale/)
- [twitter@discdiver](https://twitter.com/discdiver)
- jeffhale.medium.com/



Preflight check

See questions in Slack 😊

Planes



Flight plans

- Test flights for each plane
 - Hello world
 - Show data
 - Plot
 - Predict
 - Cruising altitude & turbulence (pros & cons)
- Grab your luggage (takeaways)
- Post-flight data (Q&A and questionnaire)
- Disembark

Flight plan materials

- GitHub repo:

<https://github.com/discdiver/dsdc-deploy-models>

Gradio

Ultralight



New, quick to fly, experimental



Gradio Demo 1: Hello world



Gradio #1: Hello world

- `pip install gradio`
- `python gradio_hello.py`

127.0.0.1:7860

PLANE

Clear

Submit

Screenshot

Flag

Gradio #1: Hello world

```
import gradio as gr

def hello(plane):
    return f"I'm an ultralight {plane} ✈️"

iface = gr.Interface(
    fn=hello,
    inputs=['text'],
    outputs=['text']
).launch()
```




Gradio Demo 2: Show data

Gradio #2: Show me the data!

```
def show_pens(alpha):  
    return pd.read_csv(  
        'https://raw.githubusercontent.com/penguins.csv')  
  
iface = gr.Interface(  
    fn=show_pens,  
    inputs=['text'],  
    outputs=[gr.outputs.DataFrame()],  
    description="Table of Palmer Penguins"  
) .launch(share=True)
```

Gradio #2: Show me the data!

- `pip install gradio pandas seaborn`
- `python gradio_pandas.py`





Gradio Demo 3: Plotting

Gradio #3: Plot it

- Plotly doesn't work as of 2.8.7 (targeted for 2.9) 😞
- You can use Matplotlib as of Gradio 2.8.2

Scatterplot of Palmer Penguins

Let's talk pens. Click to see a plot.

place holder




Clear

Submit

Output

0.3s



Flag

Talk more about Penguins here, shall we?

[view the api](#) 🔑 • built with [gradio](#) 📦

Gradio #3: Plot it

```
def plot_pens(place_holder):  
    """scatter plot penguin chars using matplotlib"""  
  
    df_pens = pd.read_csv("https://raw.githubusercontent.com/penguins.csv")  
  
    fig = plt.figure()  
    plt.scatter(  
        x=df_pens["bill_length_mm"], y=df_pens["bill_depth_mm"]  
    )  
  
    return fig
```

Gradio #3: Plot it

```
iface = gr.Interface(  
    fn=plot_pens,  
    layout="vertical",  
    inputs=["checkbox"],  
    outputs=["plot"],  
    title="Scatterplot of Palmer Penguins",  
    description="Let's talk pens. Click to see a plot.",  
    article="Talk more about Penguins here, shall we?",  
    theme="peach",  
    live=True,  
).launch()
```



Gradio Demo 4: Predict

Gradio #4: Model inference

```
import gradio as gr  
gr.Interface.load('huggingface/gpt2').launch()
```

gpt2

INPUT

Clear Submit

Screenshot Flag

Gradio #4: Predict - prettier

```
gr.Interface.load(
    "huggingface/gpt2",
    title="Storytelling with GPT2",
    css="""
        body {background: rgb(2,0,36);
            background: linear-gradient(
                180deg,
                rgba(2,0,36,1) 0%,
                rgba(7,51,99,1) 70%,
                rgba(6,3,17,1) 100%);}
        .title {color: white !important;}
    """,
).launch()
```

Gradio #4: Predict - prettier

Storytelling with GPT2

INPUT

I wish I were a fish

Clear

Submit

OUTPUT

1.68s

I wish I were a fish. I don't know if I've ever done that, but I still am a fish for good measure -- because, by the end of the day, maybe I got one back. That wasn't the case. I

Screenshot

Flag

Gradio Data API - One Click!

Response:

```
{  
  "data": [ Union[str, number] ],  
  "durations": [ float ], // the time taken for the prediction to complete  
  "avg_durations": [ float ] // the average time taken for all predictions so far (used to  
estimate the runtime)  
}
```

Try it (live demo):

Python

cURL

Javascript




```
curl -X POST http://127.0.0.1:7860/api/predict -H 'Content-Type: application/json' -d  
'{"data": [ "Hello World" ]}'
```


Gradio Pros

- Quick demos for ML 🚀
- Built-in interpretability (sometimes) 🔍
- Auto-docs 📄
- Nice Hugging Face model integration 😊
- Bright future 😎
- Easy hosting at Hugging Face spaces

<https://huggingface.co/spaces/discdiver/gpt2>

Gradio Cons

- Rough around the edges - early stage 
- Not easy to customize elements 
- Single page only 

Gradio - Hugging Face Acquisition



Abubakar Abid



Streamlit

Cessna Citation Longitude



Light, quick to takeoff, easy flying



Streamlit Demo 1: Hello world

Streamlit #1: Hello world

Hello from Streamlit!

```
import streamlit as st  
  
name = "Jeff"  
  
st.title(f"Hello from Streamlit, {name}!")
```

- *pip install streamlit*
- *streamlit run streamlit_hello.py*



Streamlit Demo 2: Show data

Streamlit #2: Show data

Streamlit with pandas

☐ Show dataframe

Streamlit #2: Show data

```
import streamlit as st

import pandas as pd

st.title("Streamlit with pandas")

show = st.checkbox("Show dataframe")

df_pens = pd.read_csv( "https://raw.githubusercontent.com/jeffhale/penguins.csv")

if show:
    df_pens
```



Streamlit Demo 3: Plotting


Streamlit #3: Plotting

Select color

☒ species

☐ island

Penguins in 3D!



species

☒ Adelie

☐ Chinstrap

☐ Gentoo

Streamlit #3: Plotting

```
choice = st.radio("Select color", ["species", "island"])
```

```
fig = px.scatter_3d(  
    data_frame=df_pens,  
    x="bill_depth_mm",  
    y="bill_length_mm",  
    z="body_mass_g",  
    color=choice,  
    title="Penguins in 3D!",  
)  
fig
```



Streamlit

Streamlit Demo 4: Predict

Streamlit #4: Predict

GPT-2 Stories

Enter your text here:

ok friends, let's talk

ok friends, let's talk

Riot - I'm on the board now. Please explain your message.

No one likes news on my channel.

Please put your email in the contact form. - Yes I have a problem

Streamlit #4: Predict

```
import streamlit as st

from transformers import pipeline

st.header("GPT-2 Stories")

input_text = st.text_area("Enter your text here:")

generator = pipeline("text-generation", model="gpt2")

output = generator(input_text, max_length=100)

output[0]["generated_text"]
```

Streamlit #4: Predict

GPT-2 Stories

Enter your text here:

ok friends, let's talk

|

ok friends, let's talk

Riot - I'm on the board now. Please explain your message.

No one likes news on my channel.

Please put your email in the contact form. - Yes I have a problem

Streamlit #4: Predict - prettier

Story time



Enter your text here:

ok smarty, here we are

I'm in another column

Here's your story:

ok smarty, here we are.

There's also the recent news of IBM's (IBM) (IBM Inc) (IBM) (IBM) (IBM MS) (IBM) smartwatches and the new BlackBerry (R) in this week's Best

Streamlit #4: Predict - prettier

```
st.header("Story time")

st.image("https://cdn.pixabay.com/photo/2017/07/12/19/03/highway-2497900_960_720.jpg")

col1, col2 = st.columns(2)

with col1:

    input_text = st.text_area("Enter your text here:")

    with st.spinner("Generating story..."):

        generator = pipeline("text-generation", model="gpt2")

        if input_text:

            generated_text = generator(input_text, max_length=60)

            st.success("Here's your story:")

            generated_text[0]["generated_text"]

with col2:

    st.header("I'm in another column")
```

Streamlit Serving Options

- Serve from Streamlit's servers for free: bit.ly/st-6plots
- Serve from Hugging Face Spaces or Heroku for free 🤗
- Pay Streamlit for more/better hosting 💰
- Host elsewhere

Streamlit Pros

- Quick websites for many Python use cases 🦆
- Many intuitive interactive widgets ✅
- Caching ⌚
- Nice hosting options 📦
- Thoughtful docs 📄
- Strong development cadence & team 💪

Streamlit Cons

- Some customizations cumbersome ✨
- Single page only (without hacks) 1

Streamlit Snowflake Acquisition

Recent acquisition by Snowflake



FastAPI

Boeing 737



Commercial grade, fast, smart!



FastAPI Demo 1: Hello world

FastAPI #1: Hello world

```
import uvicorn

from fastapi import FastAPI

app = FastAPI()

@app.get("/")
def home():
    return {"Hello world": "How's it going?"}

if __name__ == "__main__":
    uvicorn.run("fastapi_hello:app")
```

FastAPI #1: Hello world

```
pip install fastapi uvicorn
```

```
python fastapi_hello.py
```

Returns json

```
{"Hello world": "How's it going ?"}
```

FastAPI



FastAPI Speed

- Uses Starlette - ASGI (*Asynchronous Server Gateway Interface*)
- Fastest Python framework - [Techempower](#) benchmark



FastAPI Demo 2: Show data

Automatic docs

GET /form Features Form

POST /form Make Prediction

accept form submission and make prediction

Parameters

Cancel

Reset

No parameters

Request body required

application/x-www-form-urlencoded

OverallQual required
integer

OverallQual

FullBath required
integer

FullBath

GarageArea required
integer

GarageArea

LotArea required
integer

LotArea

FastAPI #2: Show me the data!

```
@app.get("/df")  
  
async def pens_data():  
    df_pens = pd.read_csv(  
        "https://raw.githubusercontent.com/jeffhale/penguins.csv"  
    )  
    df_no_nans = df_pens.fillna(-1.01)  
    return df_no_nans
```

FastAPI #2: Show me the data!

Return DataFrame, get back JSON 🕶️

```
{"species":{"0":"Adelie","1":"Adelie"...
```





FastAPI Demo 3: Plotting

FastAPI #3: Plotting



Penguins in 3D!



FastAPI #3: Plotting

```
@app.get("/plot")
async def plot() -> HTMLResponse:
    """return a plotly plot"""
    fig = px.scatter_3d(
        data_frame=df,
        x="bill_depth_mm",
        y="bill_length_mm",
        z="body_mass_g",
        color="species",
        title="Penguins in 3D!",
    )
    return HTMLResponse(fig.to_html())
```



FastAPI Demo 4: Predict

FastAPI #4: Predict: Form

Enter characteristics of your property in Ames, Iowa

Overall House Quality

Number of Full Bathrooms

Garage Area

Lot Area

FastAPI #4: Predict: Form

```
app = FastAPI()
```

```
templates = Jinja2Templates(directory="templates")
```

```
@app.get("/form", response_class=HTMLResponse)
```

```
async def features_form(request: Request):
```

```
    """form for getting data"""
```

```
    return templates.TemplateResponse(
```

```
        "form.html",
```

```
        {"request": request})
```

FastAPI #4: Predict: HTML Template

```
<body>
```

```
  <h1>Enter characteristics of your property in Ames, Iowa</h1>
```

```
  <form method='POST' enctype="multipart/form-data">
```

```
    <label for="OverallQual">Overall House Quality</label><br>
```

```
    <input type="number" name="OverallQual">
```

```
    <label for="FullBath">Number of Full Bathrooms</label><br>
```

```
    <input type="number" name="FullBath">
```

```
    <label for="GarageArea">Garage Area</label><br>
```

```
    <input type="number" name="GarageArea">
```

```
    <label for="LotArea">Lot Area</label><br>
```

```
    <input type="number" name="LotArea">
```

```
    <p><button type="submit" value="Submit">Submit</button></p>
```

```
  </form>
```

```
</body>
```

FastAPI #4: Predict: pydantic schema for form

```
class FeaturesForm(BaseModel):  
    """pydantic model to get the form input we want"""  
  
    OverallQual: int  
    FullBath: int  
    GarageArea: int  
    LotArea: int  
  
    @classmethod  
    def as_form(cls, OverallQual: int = Form(...), FullBath: int = Form(...),  
                GarageArea: int = Form(...), LotArea: int = Form(...)):  
        return cls(OverallQual=OverallQual, FullBath=FullBath,  
                    GarageArea=GarageArea, LotArea=LotArea)
```


FastAPI #4: Display prediction

Your  house is worth \$231,471.18.

Cool! 

FastAPI #4: Predict: process








```
@app.post("/form", response_class=HTMLResponse)
async def make_prediction(
    request: Request,
    user_input: FeaturesForm=Depends(FeaturesForm.as_form)
):
    """accept form submission and make prediction"""
    ...load model and make prediction

    return templates.TemplateResponse(
        "results.html", {"request": request, "prediction": pred}
    )
```

FastAPI #4: Display prediction

```
<body>
  <div class='container'>
    <div class='row'>
      <div class='col-lg text-center'>
        <h2> Your 🏠 house is worth
          {{ "${:,.2f}".format(prediction) }}.<h2>
        <h3> Cool! 🎉 </h3 </div>
      </div>
    </div>
  </body>
```

FastAPI Pros

- Fastest Python API framework - async ASGI 
- Automatic API documentation 
- Extensive docs 
- Nice test client 
- Nice dependency injection 
- Data validation with pydantic / SQL Model integration 
- Security / authentication support 

FastAPI Pros




Sebastian Ramirez


FastAPI Cons

- Reliant on a single maintainer 🤖
- Overriding uvicorn logging is a bit of a pain 🌳
- HTML templating more painful than Flask ⌨️

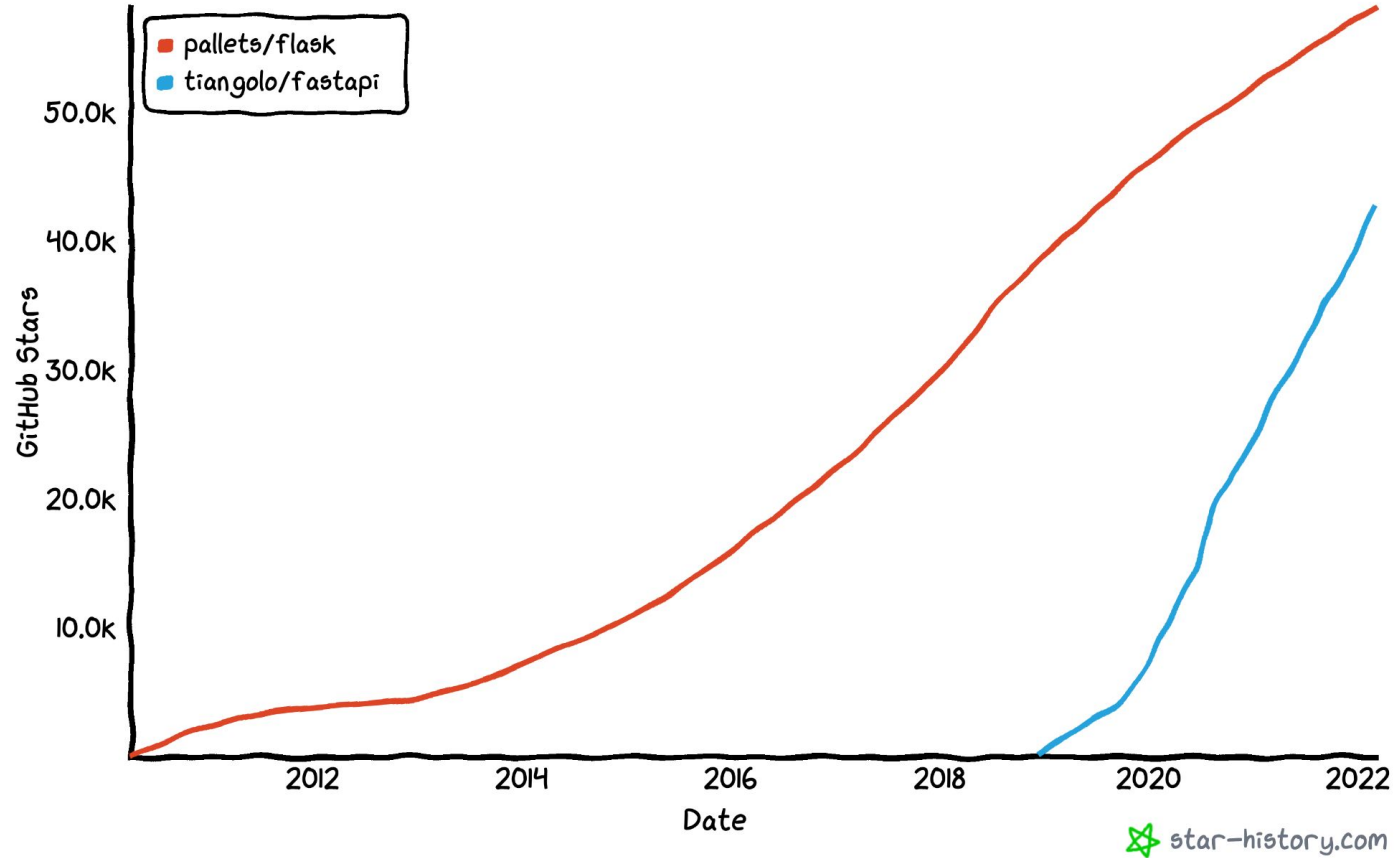
Grab your luggage (take aways)

	Web App	Data API
	Yes	Yes
 Streamlit	Yes	No
 FastAPI	Yes (Jinja templates)	Yes

Star history



Star history



star-history.com

What about flask?

- Huge inspiration for FastAPI
- Now has more async support, but not easily full async
- FastAPI is faster
- FastAPI leverages typing
- FastAPI winning mindshare
- Gradio just switched to FastAPI backend
- Flask is nicer for quick web app

**Grab Your Luggage
(takeaways)**

**Use what you know, unless it
doesn't meet your needs**

Blank slate?



**Learn what's popular, growing,
and quick to get off the ground**



Streamlit



**For single-page app that doesn't
need custom styling**



Gradio for quick 🤗 models for fun



FastAPI for serving data

What to learn next?

(Newer to Python)



Streamlit

Post-flight

- Questions?
- Please fill out short form:

<https://forms.gle/pEz4suEqZTbBHhdC7>

Disembark

Thank you for flying the deployment skies!

Jeff Hale


- [linkedin.com/in/-jeffhale/](https://www.linkedin.com/in/-jeffhale/)
- [@discdiver](#)
- jeffhale.medium.com/



Versions used

- Gradio 2.8.7
- Streamlit 1.5.1
- FastAPI 0.75.0

Python Tools to Deploy Your Machine Learning Models Faster 🚀



Jeff Hale

